

AI PROJECT PLAN

Predicting Student Performance

Using Machine Learning



Yuva Internship | Artificial Intelligence Trainee

Week 1 Submission | April 04, 2026

Executive Summary: This document presents a comprehensive project plan for an AI-driven predictive analytics initiative aimed at forecasting student academic performance. Leveraging publicly available educational datasets, the project proposes a full machine-learning pipeline — from data ingestion and feature engineering through model training, evaluation, and ethical deployment — with a structured 10-week timeline and clearly defined milestones.

Table of Contents

1. Project Overview and Objectives
2. Hypothesis Formulation
3. Tools and Technology Selection
4. Methodology: Data Collection and Analysis
5. Model Training Strategy

6. Evaluation Techniques
7. Potential Challenges
8. Risk Mitigation Strategies
9. Data Governance and Scalability
10. Ethical Considerations
11. Expected Outcomes
12. Project Timeline and Gantt Chart

1. Project Overview and Objectives

1.1 Background

Education is one of the most critical sectors where data-driven decision-making can create a profound impact. Academic institutions worldwide are increasingly generating large volumes of student-related data — from attendance records and assessment scores to socioeconomic indicators and learning behaviour logs. However, much of this data remains underutilised. Early identification of at-risk students can allow educators and administrators to intervene before a student falls too far behind, significantly improving learning outcomes and institutional performance metrics.

This project proposes the development of a machine learning (ML) system that leverages publicly available educational datasets to predict student performance — specifically, end-of-term grades and the likelihood of course failure. The system will be designed to be interpretable, fair, and deployable within real educational environments.

1.2 Project Objectives

- Build a supervised ML pipeline that accurately predicts student academic performance using historical and demographic data.
- Identify the most significant features (e.g., study time, parental education, prior grades) influencing academic outcomes.
- Develop an early-warning module to flag at-risk students before final assessments.
- Ensure the model is interpretable by educators using explainable AI (XAI) techniques such as SHAP.
- Document all design decisions, assumptions, and ethical considerations to support responsible deployment.
- Produce a scalable architecture that can be extended to other educational institutions or subjects.

1.3 Problem Statement

Problem: A significant proportion of students underperform or drop out each academic year. Educators often lack timely, data-driven tools to identify such students early enough to intervene effectively. This project addresses that gap by building an ML model that provides actionable, predictive insights from readily available student data.

2. Hypothesis Formulation

Before embarking on model development, it is essential to define working hypotheses that guide feature selection, model choice, and interpretation. These hypotheses are grounded in existing educational research and will be tested empirically during the data analysis phase.

- **H1 – Study Time:** Students who dedicate more weekly hours to studying achieve significantly higher final grades. We expect a positive, non-linear correlation between study time and performance.
- **H2 – Prior Academic Performance:** A student's grade in the first and second assessment periods (G1, G2) is the strongest single predictor of their final grade (G3). Prior performance encodes learning ability, consistency, and engagement simultaneously.
- **H3 – Parental Education:** Higher parental education levels correlate positively with student performance due to access to academic support, resources, and motivational factors at home.
- **H4 – Absenteeism:** Higher levels of school absence negatively impact academic performance through a direct causal mechanism — missed instruction leads to knowledge gaps.
- **H5 – Internet Access:** Students with home internet access will, on average, perform better due to greater access to supplementary learning materials and online resources.
- **H6 – Extracurricular Activities:** Moderate participation in extracurricular activities has a positive effect on performance, while excessive participation may be detrimental.

These hypotheses will be validated through Exploratory Data Analysis (EDA), correlation matrices, and feature importance scores derived from the trained models.

3. Tools and Technology Selection

3.1 Programming Environment

The entire project will be developed in Python 3.10+, chosen for its extensive ML ecosystem, readability, and wide adoption in both academia and industry.

3.2 Core Libraries and Frameworks

Category	Tool / Library	Purpose
Data Handling	Pandas, NumPy	Data loading, cleaning, and manipulation
Visualisation	Matplotlib, Seaborn, Plotly	EDA charts, correlation heatmaps, feature plots
ML Framework	Scikit-learn	Model training, cross-validation, pipelines
Boosting Models	XGBoost, LightGBM	High-performance gradient boosting classifiers
Deep Learning	TensorFlow / Keras	Optional neural network baseline
Explainability	SHAP, LIME	Model interpretability and feature attribution
Experiment Mgmt	MLflow	Experiment tracking, model versioning
Notebook Env.	Jupyter Lab	Interactive development and documentation

Category	Tool / Library	Purpose
Version Control	Git / GitHub	Code versioning and collaboration
Deployment	Flask / Streamlit	Web-based demo dashboard for educators

Table 1: Technology Stack Overview

3.3 Dataset

The primary dataset is the **UCI Student Performance Dataset** (Cortez & Silva, 2008), freely available on the UCI Machine Learning Repository and Kaggle. It contains 649 records across 33 attributes for secondary school students in Portugal, covering demographics, social factors, school-related variables, and three periodic grades (G1, G2, G3). A secondary dataset from the **Open University Learning Analytics Dataset (OULAD)** may be incorporated for multi-institutional validation.

4. Methodology: Data Collection and Analysis

4.1 Data Collection

Data will be sourced from publicly available repositories without any privacy concerns. The UCI Student Performance Dataset will be downloaded programmatically via the UCI ML Repository API. If additional data is required, synthetic data augmentation techniques (SMOTE for class imbalance) will be applied to enrich the training set.

4.2 Data Preprocessing Pipeline

- **Data Loading:** Ingest raw CSV files into Pandas DataFrames.
- **Missing Value Treatment:** Impute numerical nulls with median; categorical nulls with mode. Columns with >30% missing data will be dropped.
- **Outlier Detection:** Use IQR method and Z-score analysis to identify and cap outliers.
- **Encoding:** Apply Label Encoding for ordinal variables (e.g., education level) and One-Hot Encoding for nominal variables (e.g., school, guardian).
- **Feature Scaling:** Apply StandardScaler for linear models; tree-based models will use raw values.
- **Feature Engineering:** Create composite features — e.g., 'average_grade' = mean(G1, G2); 'support_score' combining family and school support indicators.
- **Train-Test Split:** 80% training / 20% testing with stratification on the target variable.

4.3 Exploratory Data Analysis (EDA)

A thorough EDA will be conducted prior to modelling. This includes: univariate distribution analysis of all features, bivariate correlation analysis (Pearson and Spearman coefficients), visualisation of grade

distributions across demographic groups, and a heatmap of the full correlation matrix. EDA findings will directly inform feature selection and hypothesis validation.

4.4 Target Variable Definition

The target variable will be defined in two ways: (a) **Regression task** — predicting the continuous final grade G3 (scale 0–20); and (b) **Classification task** — predicting a binary label where $G3 < 10$ = 'At Risk' (1) and $G3 \geq 10$ = 'Passing' (0). Both formulations will be explored and compared.

5. Model Training Strategy

A multi-model approach will be adopted to identify the best-performing algorithm for this dataset. The following models will be trained and compared:

Category	Algorithm	Rationale
Baseline	Logistic Regression / Linear Regression	Establishes a simple, interpretable benchmark.
Tree-Based	Decision Tree, Random Forest	Handles non-linearity; Random Forest reduces overfitting via ensemble averaging.
Boosting	XGBoost, LightGBM	State-of-the-art performance on tabular data with regularisation built in.
SVM	Support Vector Machine (RBF kernel)	Effective in high-dimensional spaces; useful for the classification task.
Neural Net	MLP (Keras, 3 hidden layers)	Optional deep learning baseline to assess gains over classical methods.

Table 2: Model Selection Summary

5.1 Hyperparameter Tuning

GridSearchCV and RandomizedSearchCV from scikit-learn will be used for hyperparameter optimisation. For XGBoost and LightGBM, Bayesian Optimisation via the Optuna library will be explored for efficiency. All tuning will be performed within a cross-validation framework to prevent data leakage.

5.2 Cross-Validation Strategy

Stratified K-Fold cross-validation (K=5) will be applied to all models during training. This ensures balanced representation of the target class across all folds, providing robust and unbiased performance estimates. A separate holdout test set (20%) will only be accessed once at the end for final evaluation.

6. Evaluation Techniques

6.1 Classification Metrics

6.2 Regression Metrics

Metric	Purpose
Accuracy	Overall correctness; reported alongside other metrics due to class imbalance sensitivity.
Precision & Recall	Critical for the 'At Risk' class — high recall minimises missed interventions.
F1-Score	Harmonic mean of Precision and Recall; primary metric for classification.
ROC-AUC	Measures discriminative ability across all classification thresholds.
Confusion Matrix	Visual breakdown of true/false positives and negatives.
MAE (Mean Absolute Error)	Average magnitude of prediction error in grade units.
RMSE (Root Mean Squared Error)	Penalises large errors more heavily; important for grade prediction.
R-squared (R2)	Proportion of variance explained by the model.

Table 3: Evaluation Metrics

6.3 Explainability

Beyond numeric metrics, SHAP (SHapley Additive exPlanations) values will be computed for the best-performing model to visualise individual feature contributions. Summary SHAP plots and force plots will be generated to help educators understand why a specific student was flagged as at-risk — a critical requirement for trust and adoption in real classrooms.

7. Potential Challenges

- Small Dataset Size:** The UCI dataset contains only ~649 records, which may limit model generalisation and statistical power for complex models.
- Class Imbalance:** The proportion of failing students may be significantly lower than passing students, causing models to be biased toward the majority class.
- Feature Overlap:** G1 and G2 are highly correlated with G3, which may overshadow the predictive power of other important socioeconomic features.
- Multicollinearity:** Several demographic and social features may be intercorrelated, complicating coefficient interpretation in linear models.
- Generalisability:** Data from Portuguese schools may not generalise well to different cultural, economic, or educational contexts.
- Overfitting Risk:** Complex models like neural networks and gradient boosting may overfit the small dataset without careful regularisation.
- Missing Context:** The dataset lacks real-time behavioural data (e.g., LMS activity, assignment submission timing) which could further improve predictions.

Ethical Sensitivity: Predictions based on socioeconomic attributes (parental education, family income) risk perpetuating systemic inequalities if misused.

8. Risk Mitigation Strategies

Risk	Likelihood	Impact	Mitigation Strategy
Small dataset	High	Medium	Apply SMOTE for augmentation; use cross-validation; consider transfer learning.
Class imbalance	High	High	Use SMOTE, class_weight='balanced'; evaluate with F1 and AUC rather than accuracy.
Overfitting	Medium	High	Apply regularisation (L1/L2), early stopping, dropout; use 5-fold CV.
Feature leakage	Medium	Critical	Strictly separate train/test sets; apply preprocessing inside CV pipeline.
Model bias	Medium	High	Conduct fairness audits; remove protected attributes from final model if bias detected.
Generalisation failure	Low	Medium	Validate on OULAD dataset; document scope and limitations clearly.
Tool incompatibility	Low	Low	Use requirements.txt and virtual environments; containerise with Docker.

Table 4: Risk Register and Mitigation Strategies

9. Data Governance and Scalability

9.1 Data Governance

Although the dataset is publicly available and anonymised, robust data governance practices will be maintained throughout the project:

- All data will be stored in a secure local environment; no personally identifiable information (PII) will be used.
- Data provenance will be documented — source, version, date of access, and any transformations applied.
- A data dictionary will be maintained describing each feature, its type, and its origin.
- Any synthetic data generation will be clearly flagged and separated from original data.
- The project will comply with principles of GDPR and India's DPDP Act where applicable.

9.2 Scalability Considerations

The architecture is designed with scalability in mind to support future expansion:

- **Modular Pipeline:** Each stage (ingestion, preprocessing, training, evaluation) is encapsulated as an independent module for easy replacement or upgrading.
- **Cloud Readiness:** The pipeline can be deployed on AWS SageMaker, Google Cloud AI Platform, or Azure ML with minimal refactoring.

- **Database Integration:** A relational database (PostgreSQL) or a data warehouse (BigQuery) can replace flat CSV files for production-scale data.
- **API Layer:** A REST API (FastAPI) will wrap the trained model, enabling integration with Student Information Systems (SIS).
- **Retraining:** The system will support scheduled retraining as new student cohort data becomes available each semester.

10. Ethical Considerations

Ethical AI development is not an afterthought but a foundational design principle of this project. The following ethical dimensions have been identified and addressed:

Fairness and Non-Discrimination: The model will be tested for disparate impact across gender, parental education level, and urban/rural background. Features that are proxies for protected attributes will be carefully evaluated and may be excluded from the final model.

Transparency and Explainability: All predictions will be accompanied by SHAP-based explanations. Educators will be provided with human-readable rationale for each at-risk flag, enabling informed and empathetic intervention.

Accountability: The model is a decision-support tool, not a decision-making tool. Final decisions regarding student support must remain with qualified educators. A clear accountability chain will be documented.

Privacy by Design: Even though the dataset is public, the system will be built assuming real-world deployment where student data is sensitive. Differential privacy techniques will be explored for future iterations.

Avoiding Labelling Harm: Labelling students as 'at risk' carries potential stigma. Communication of model outputs to students and parents must be handled sensitively, framed as support opportunities rather than predictions of failure.

Consent and Autonomy: In any real deployment, informed consent must be obtained from students and guardians before their data is used in predictive systems.

11. Expected Outcomes

Upon successful completion of this project, the following deliverables and outcomes are anticipated:

Deliverable	Description
Predictive Model	A trained and validated ML model achieving at least 85% accuracy ($F1 \geq 0.82$) on the classification task.

Deliverable	Description
Feature Importance Report	A ranked list of the top 10 features most predictive of student performance, validated across multiple cohorts.
Early-Warning Dashboard	A Streamlit-based web application that allows educators to input student data and receive real-time risk alerts.
Technical Documentation	Full codebase with docstrings, a reproducible Jupyter Notebook, and an MLflow experiment log.
Research Summary	A 2-page summary of findings suitable for sharing with academic stakeholders or for submission to a journal.
Ethical Review Report	A fairness audit report documenting model performance across demographic subgroups and a plan for mitigation.

Table 5: Expected Project Deliverables

11.1 Impact Assessment

If deployed in a real educational environment, this system is expected to reduce the rate of undetected at-risk students by 30–40%, allowing for timely academic interventions. For the institution, this translates to improved pass rates, better student retention, and data-driven resource allocation. The project also serves as a proof-of-concept for responsible AI adoption in education — a sector where trust, fairness, and interpretability are paramount.

12. Project Timeline and Gantt Chart

The project is planned over a 10-week cycle from kickoff to final reporting. The timeline below outlines all major phases and deliverables.

Phase / Task	W1	W2	W3	W4	W5	W6	W7	W8	W9	W10
1. Project Setup & Literature Review	■	■								
2. Data Acquisition & EDA		■	■							
3. Data Preprocessing Pipeline			■	■						
4. Baseline Model Training				■	■					
5. Advanced Model Training & Tuning					■	■	■			
6. Model Evaluation & SHAP Analysis							■	■		
7. Dashboard Development								■	■	
8. Ethical Audit & Fairness Testing							■	■	■	
9. Documentation & Reporting									■	■
10. Final Review & Submission										■

Figure 1: Project Gantt Chart (W = Week)

12.1 Key Milestones

Milestone Date	Deliverable
End of Week 2	Literature review complete; dataset downloaded and explored.
End of Week 4	Clean dataset and feature-engineered training set ready.
End of Week 5	Baseline model trained with initial benchmark metrics.
End of Week 7	Best model selected after hyperparameter tuning and cross-validation.
End of Week 8	SHAP explainability report and model evaluation complete.
End of Week 9	Working dashboard prototype and fairness audit finished.
End of Week 10	Final documentation, codebase, and submission ready.

Table 6: Key Project Milestones

References

- [1] Cortez, P., & Silva, A. (2008). Using data mining to predict secondary school student performance. *Proceedings of the 5th Annual Future Business Technology Conference (FUBUTEC 2008)*, Porto, Portugal. Available: <https://archive.ics.uci.edu/dataset/320/student+performance>
- [2] Romero, C., & Ventura, S. (2010). Educational data mining: A review of the state of the art. *IEEE Transactions on Systems, Man, and Cybernetics*, 40(6), 601–618.
- [3] Lundberg, S. M., & Lee, S.-I. (2017). A unified approach to interpreting model predictions. *Advances in Neural Information Processing Systems*, 30.
- [4] Open University Learning Analytics Dataset (OULAD). (2017). The Open University. Available: https://analyse.kmi.open.ac.uk/open_dataset
- [5] Chawla, N. V., Bowyer, K. W., Hall, L. O., & Kegelmeyer, W. P. (2002). SMOTE: Synthetic minority over-sampling technique. *Journal of Artificial Intelligence Research*, 16, 321–357.
- [6] Pedregosa, F., et al. (2011). Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12, 2825–2830.
- [7] Chen, T., & Guestrin, C. (2016). XGBoost: A scalable tree boosting system. *Proceedings of the 22nd ACM SIGKDD*, San Francisco, CA.
- [8] Zawacki-Richter, O., et al. (2019). Systematic review of research on artificial intelligence applications in higher education. *International Journal of Educational Technology in Higher Education*, 16(39).